



Chap. 3) Processes

경희대학교 컴퓨터공학과

조진성

Program vs. Process

■ Program

- ✓ Executable file on a disk
- ✓ Loaded into memory and executed by the kernel

■ Process

- ✓ Executing instance of a program
- ✓ The basic unit of execution and scheduling
- ✓ A process is named using its process ID (PID)



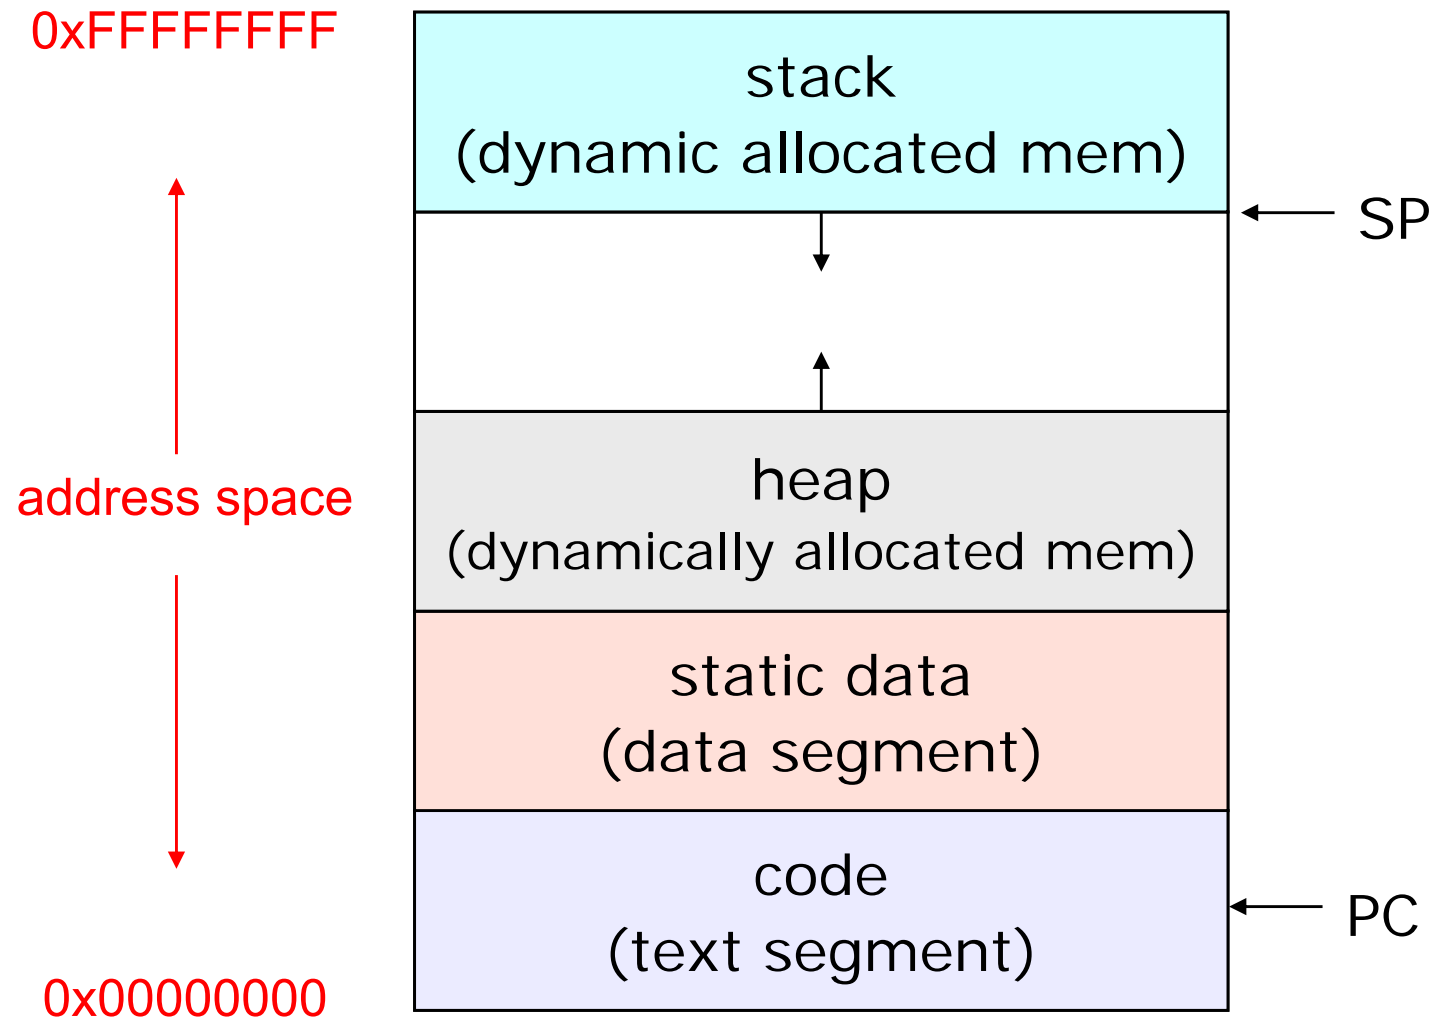
Process Concept

■ What is the process?

- ✓ An instance of a program in execution
- ✓ An encapsulation of the flow of control in a program
- ✓ A dynamic and active entity
- ✓ The basic unit of execution and scheduling
- ✓ A process is named using its process ID (PID)

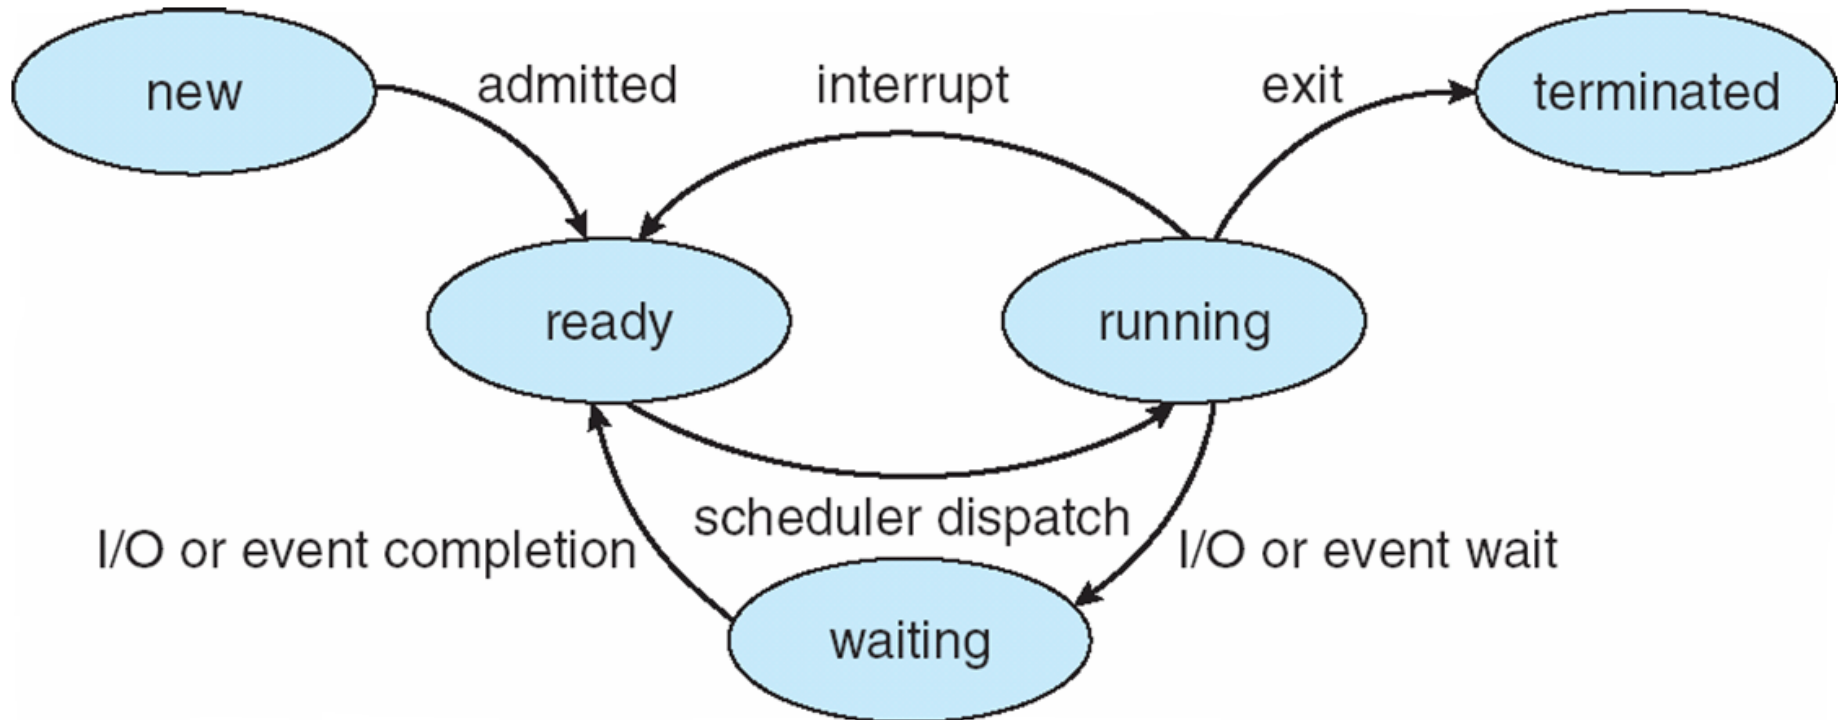


Process Address Space



Process State

■ State diagram



Process State Transition

Linux example

```
xterm
339 ?      S      0:21  clanmgr
340 ?      S      0:00  clanmgr
345 ?      S      0:00  clanagent
346 ?      S      0:00  clanagent
596 ?      S      0:00  syslogd -m 0
601 ?      S      0:00  klogd -x
621 ?      S      0:00  portmap
649 ?      S      0:00  rpc.statd
761 ?      S      0:00  /usr/sbin/apmd -p 10 -w 5 -W -P /etc/sysconfig/apm-sc
821 ?      S      0:00  /usr/sbin/automount --timeout 300 /user file /etc/aut
843 ?      S      0:00  /usr/sbin/sshd
863 ?      S      0:00  xinetd -stayalive -reuse -pidfile /var/run/xinetd.pid
905 ?      S      0:00  sendmail: accepting connections
924 ?      S      0:01  gpm -t imps2 -m /dev/mouse
942 ?      S      0:00  crond
1016 ?     S      0:00  xfs -droppriv -daemon
1052 ?     S      0:00  /usr/sbin/atd
1059 ?     S      0:00  login -- jinsoo
1060 tty2   S      0:00  /sbin/mingetty tty2
1061 tty3   S      0:00  /sbin/mingetty tty3
1062 tty4   S      0:00  /sbin/mingetty tty4
1063 tty5   S      0:00  /sbin/mingetty tty5
1064 tty6   S      0:00  /sbin/mingetty tty6
27499 ?    SW     0:00  [rpciod]
27500 ?    SW     0:00  [lockd]
27501 tty1  S      0:00  -tcsh
5365 ?     S      0:00  /usr/sbin/sshd
5367 pts/0  S      0:00  -tcsh
5394 pts/0  R      0:00  ps ax
[oz0:/-5]
```

- R: Runnable
- S: Sleeping
- T: Traced or Stopped
- D: Uninterruptible Sleep
- Z: Zombie
- W: No resident pages
- <: High-priority task
- N: Low-priority task
- L: Has pages locked into memory



Process Control Block (PCB)

■ Information associated with each process

- ✓ Process state
- ✓ Program counter
- ✓ CPU registers
- ✓ CPU scheduling information
- ✓ Memory-management information
- ✓ Accounting information
- ✓ I/O status information

■ Cf) *task_struct* in Linux

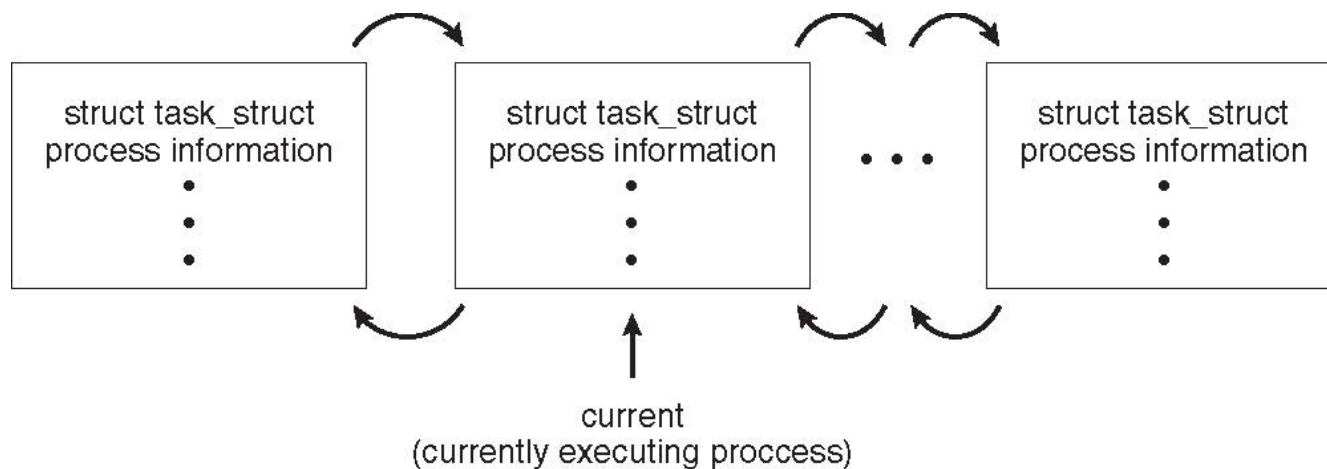
- ✓ 1456 bytes as of Linux 2.4.18



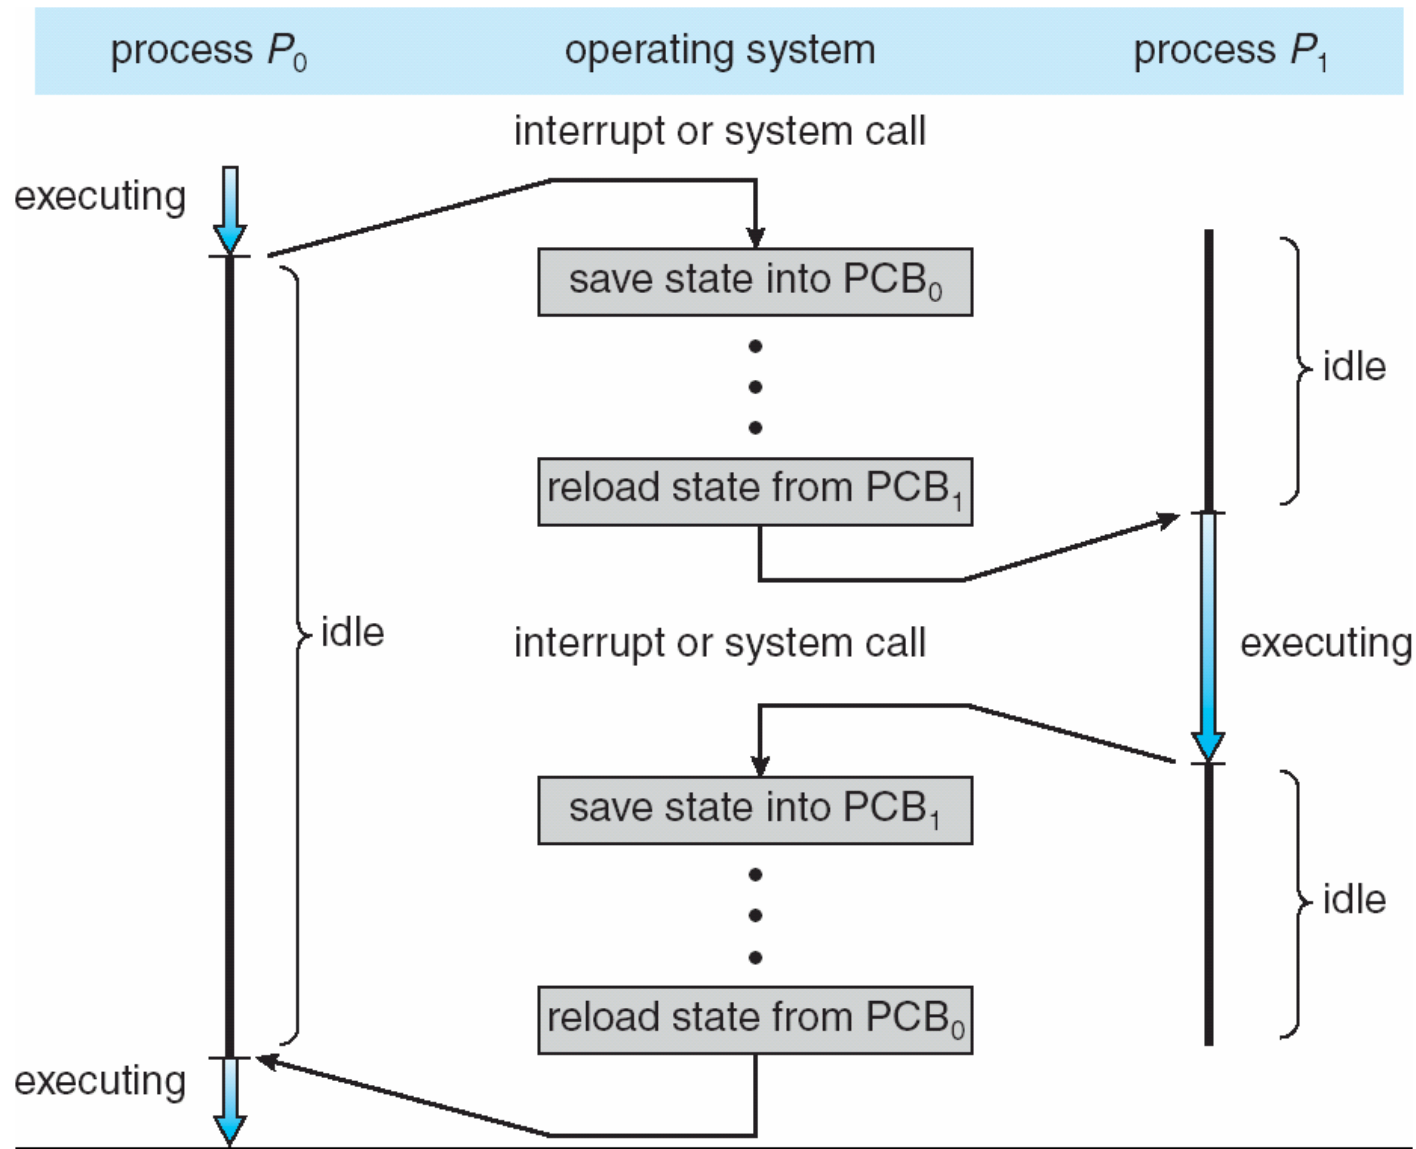
Process Control Block (PCB)

Represented by the C structure `task_struct`

```
pid t_pid; /* process identifier */
long state; /* state of the process */
unsigned int time_slice /* scheduling information */
struct task_struct *parent; /* this process's parent */
struct list_head children; /* this process's children */
struct files_struct *files; /* list of open files */
struct mm_struct *mm; /* address space of this process */
```



Context Switch (CPU Switch)



Context Switch

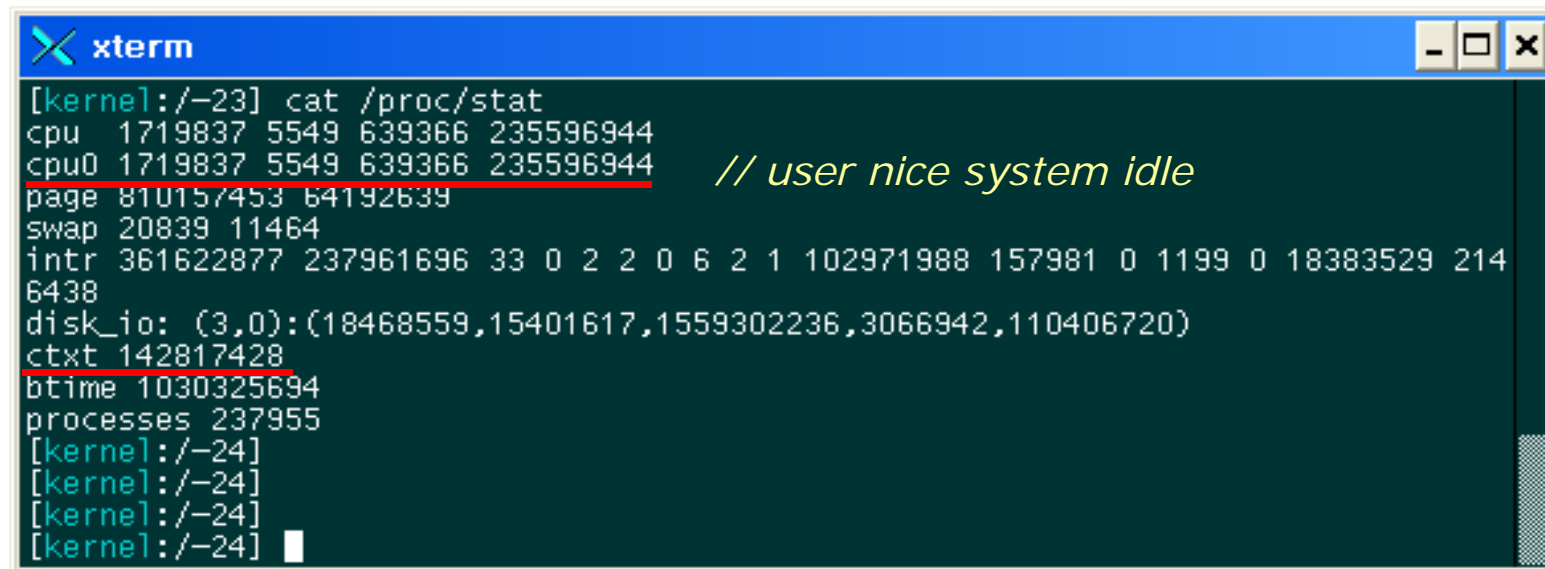
- The act of switching the CPU from one process to another
- Administrative overhead
 - ✓ saving and loading registers and memory maps
 - ✓ flushing and reloading the memory cache
 - ✓ updating various tables and lists, etc.
- Context switch overhead is dependent on hardware support
 - ✓ Multiple register sets in UltraSPARC
 - ✓ Advanced memory management techniques may require extra data to be switched with each context
- 100s or 1000s of switches/s typically



Context Switch

■ Linux example

- ✓ Total 237,961,696 ticks = 661 hours = 27.5 days
- ✓ Total 142,817,428 context switches
- ✓ Roughly 60 context switches / sec



```
xterm
[kernel]:/-23] cat /proc/stat
cpu 1719837 5549 639366 235596944
cpu0 1719837 5549 639366 235596944 // user nice system idle
page 810157453 64192639
swap 20839 11464
intr 361622877 237961696 33 0 2 2 0 6 2 1 102971988 157981 0 1199 0 18383529 214
6438
disk_io: (3,0):(18468559,15401617,1559302236,3066942,110406720)
ctxt 142817428
btime 1030325694
processes 237955
[kernel]:/-24]
[kernel]:/-24]
[kernel]:/-24]
[kernel]:/-24]
```



Schedulers

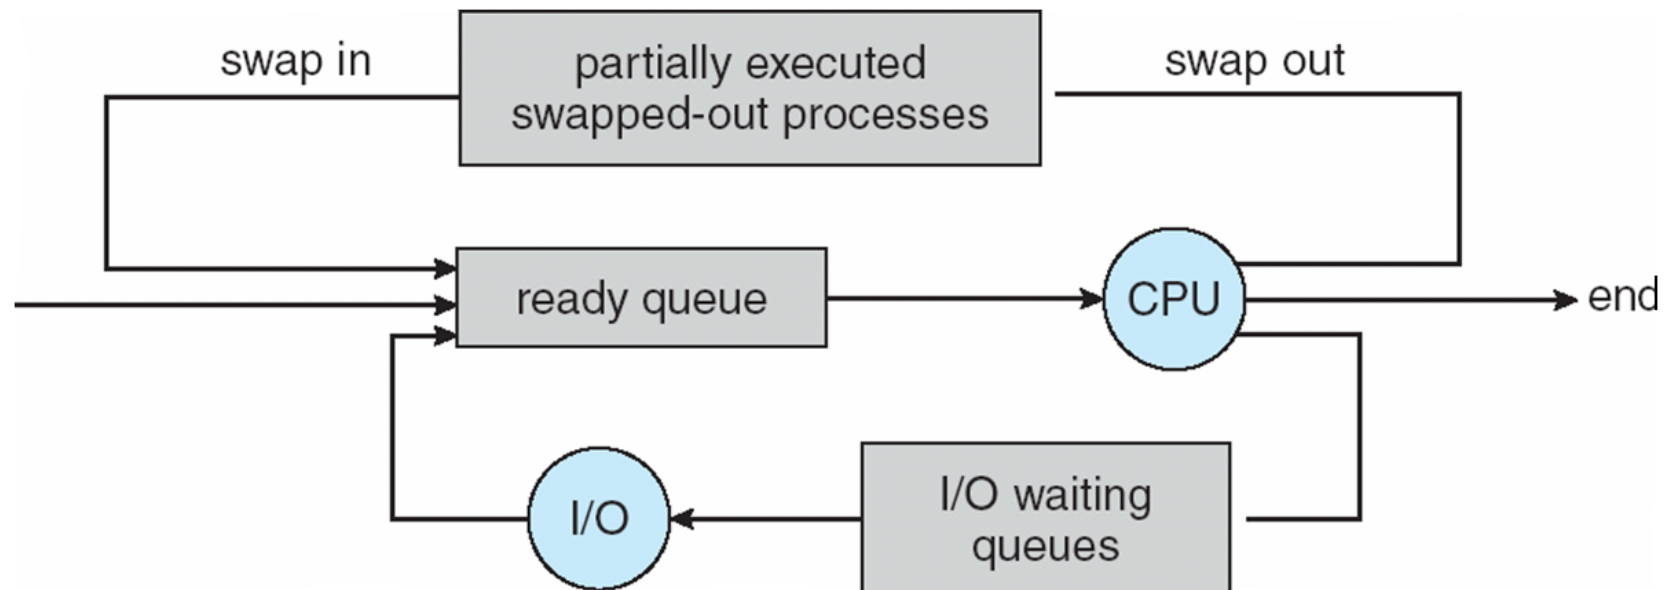
- Long-term scheduler (or job scheduler)

- ✓ selects which processes should be brought into the ready queue

- Short-term scheduler (or CPU scheduler)

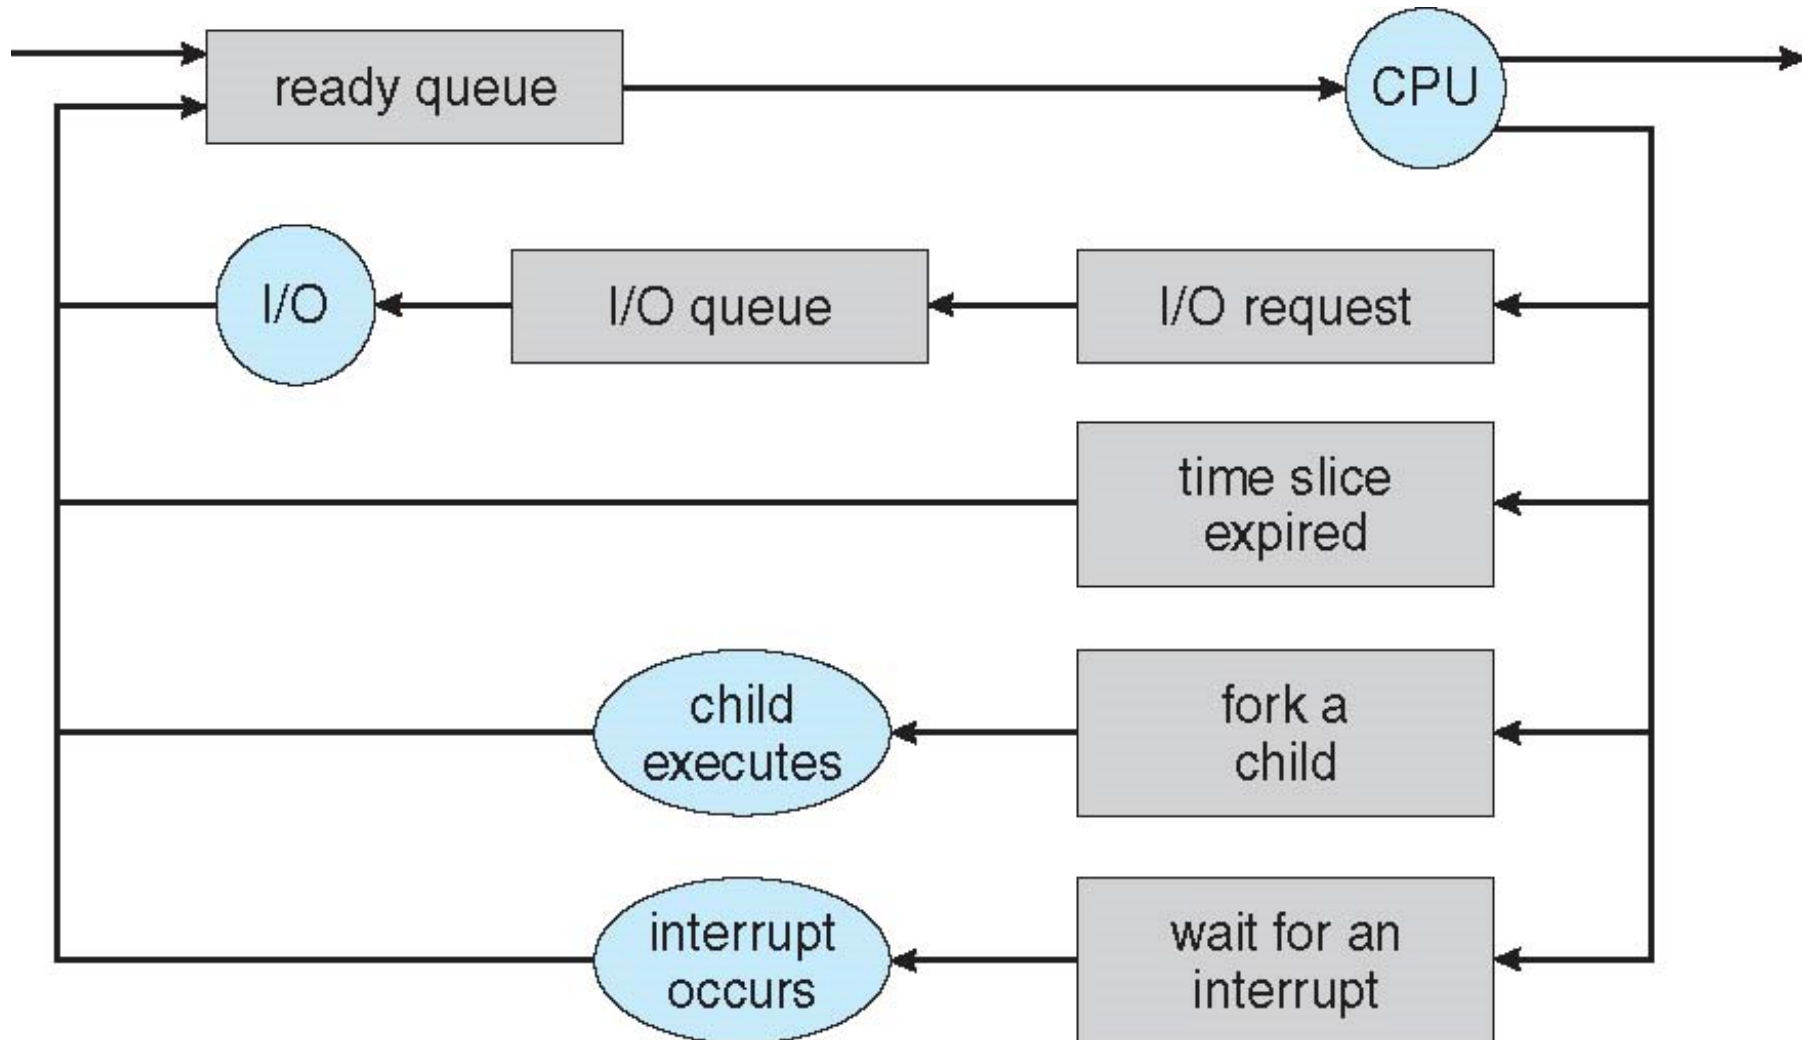
- ✓ selects which process should be executed next and allocates CPU

- Medium-term scheduler (or swapper)



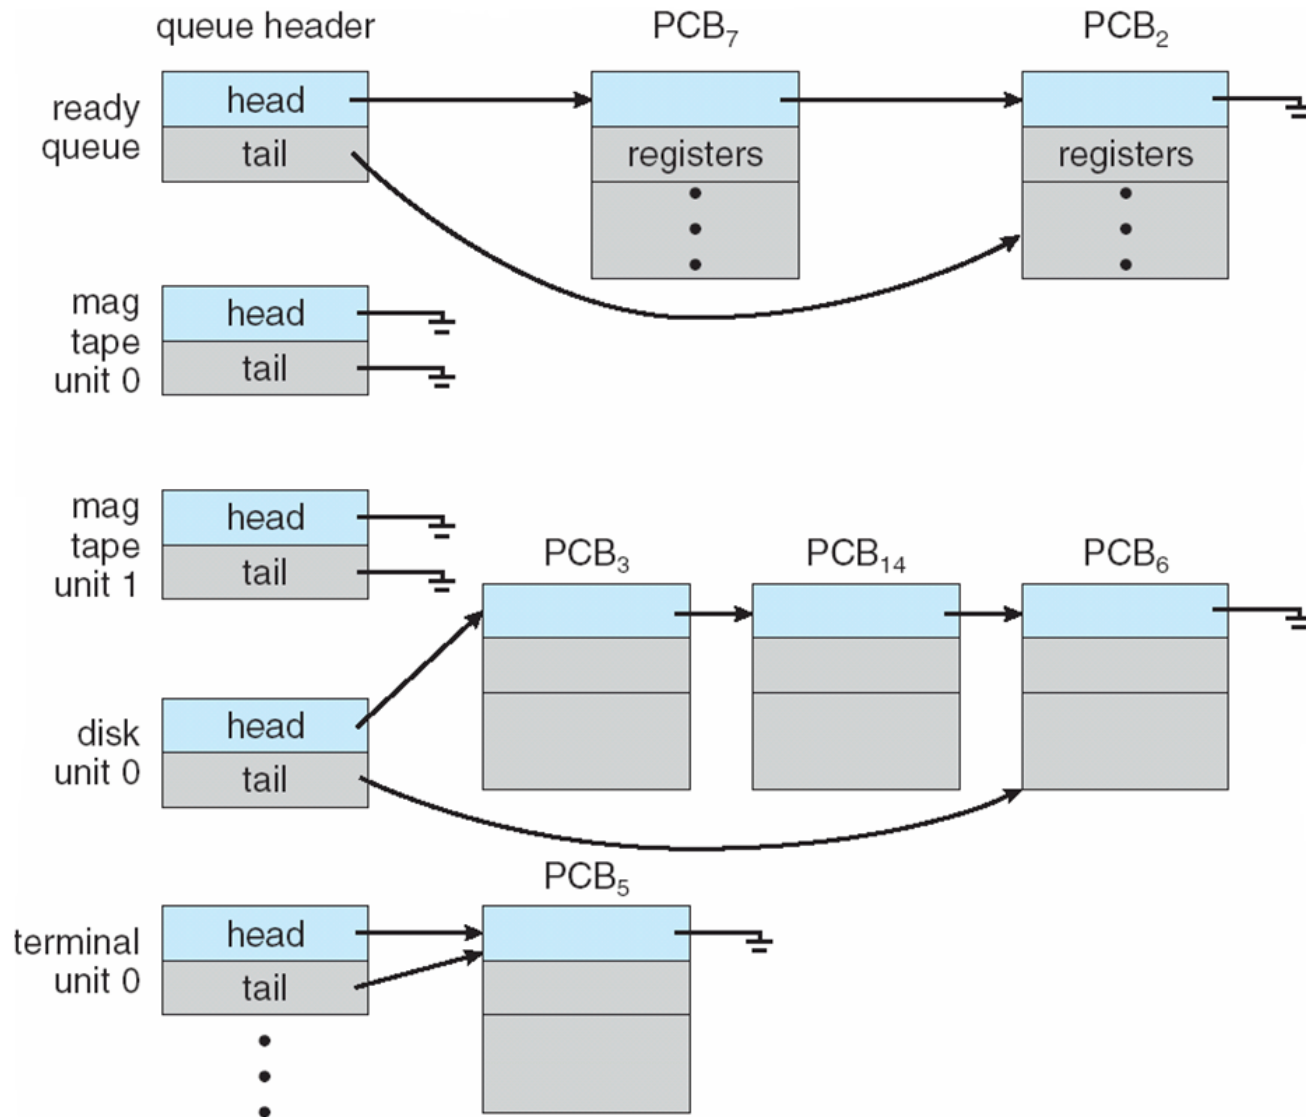
Representation of Process Scheduling

■ Queueing diagram



Representation of Process Scheduling

■ Ready queue and various I/O device queues



Operations on Processes

■ Process creation

- ✓ `fork()`

■ Process execution

- ✓ `exec()`

■ Process termination

- ✓ `exit()`
- ✓ `_exit()`
- ✓ `abort()`
- ✓ `wait()`

■ Cooperating processes

- ✓ Inter-Process Communication (IPC)



Process Creation: Unix/Linux

```
int fork()
```

■ fork()

- ✓ Creates and initializes a new PCB
- ✓ Creates and initializes a new address space
- ✓ Initializes the address space with a copy of the entire contents of the address space of the parent
- ✓ Initializes the kernel resources to point to the resources used by parent (e.g., open files)
- ✓ Places the PCB on the ready queue
- ✓ Returns the child's PID to the parent, and zero to the child



Process Creation: Unix/Linux

■ Sharing of open files between parent and child after `fork`

Parent Process Table

	fd flags	ptr
fd 0 :		
fd 1 :		
fd 2 :		

Child Process Table

	fd flags	ptr
fd 0 :		
fd 1 :		
fd 2 :		

File Table

File status flags
Current file offset
v-node ptr

File status flags
Current file offset
v-node ptr

File status flags
Current file offset
v-node ptr

v-node Table

v-node information
i-node information

v-node information
i-node information

v-node information
i-node information



fork()

```
#include <sys/types.h>
#include <unistd.h>

int main()
{
    int pid;

    if ((pid = fork()) == 0)
        /* child */
        printf ("Child of %d is %d\n", getppid(), getpid());
    else
        /* parent */
        printf ("I am %d. My child is %d\n", getpid(), pid);
}
```



fork(): Example Output

% ./a.out

I am 31098. My child is 31099.

Child of 31098 is 31099.

% ./a.out

Child of 31100 is 31101.

I am 31100. My child is 31101.



Why fork()?

- Very useful when the child...
 - ✓ is cooperating with the parent
 - ✓ relies upon the parent's data to accomplish its task
 - ✓ Example: Web server

```
While (1) {  
    int sock = accept();  
    if ((pid = fork()) == 0) {  
        /* Handle client request */  
    } else {  
        /* Close socket */  
    }  
}
```

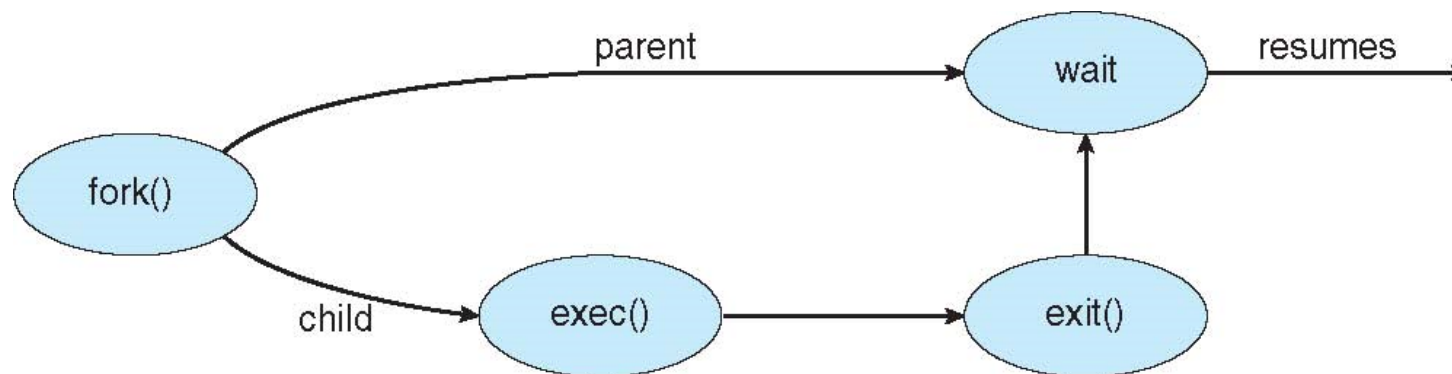


Process Execution: Unix/Linux

```
int exec (char *prog, char *argv[])
```

■ exec()

- ✓ Stops the current process
- ✓ Loads the program “prog” into the process’ address space
- ✓ Initializes hardware context and args for the new program
- ✓ Places the PCB on the ready queue
 - Note: exec() does not create a new process
- ✓ What does it mean for exec() to return?

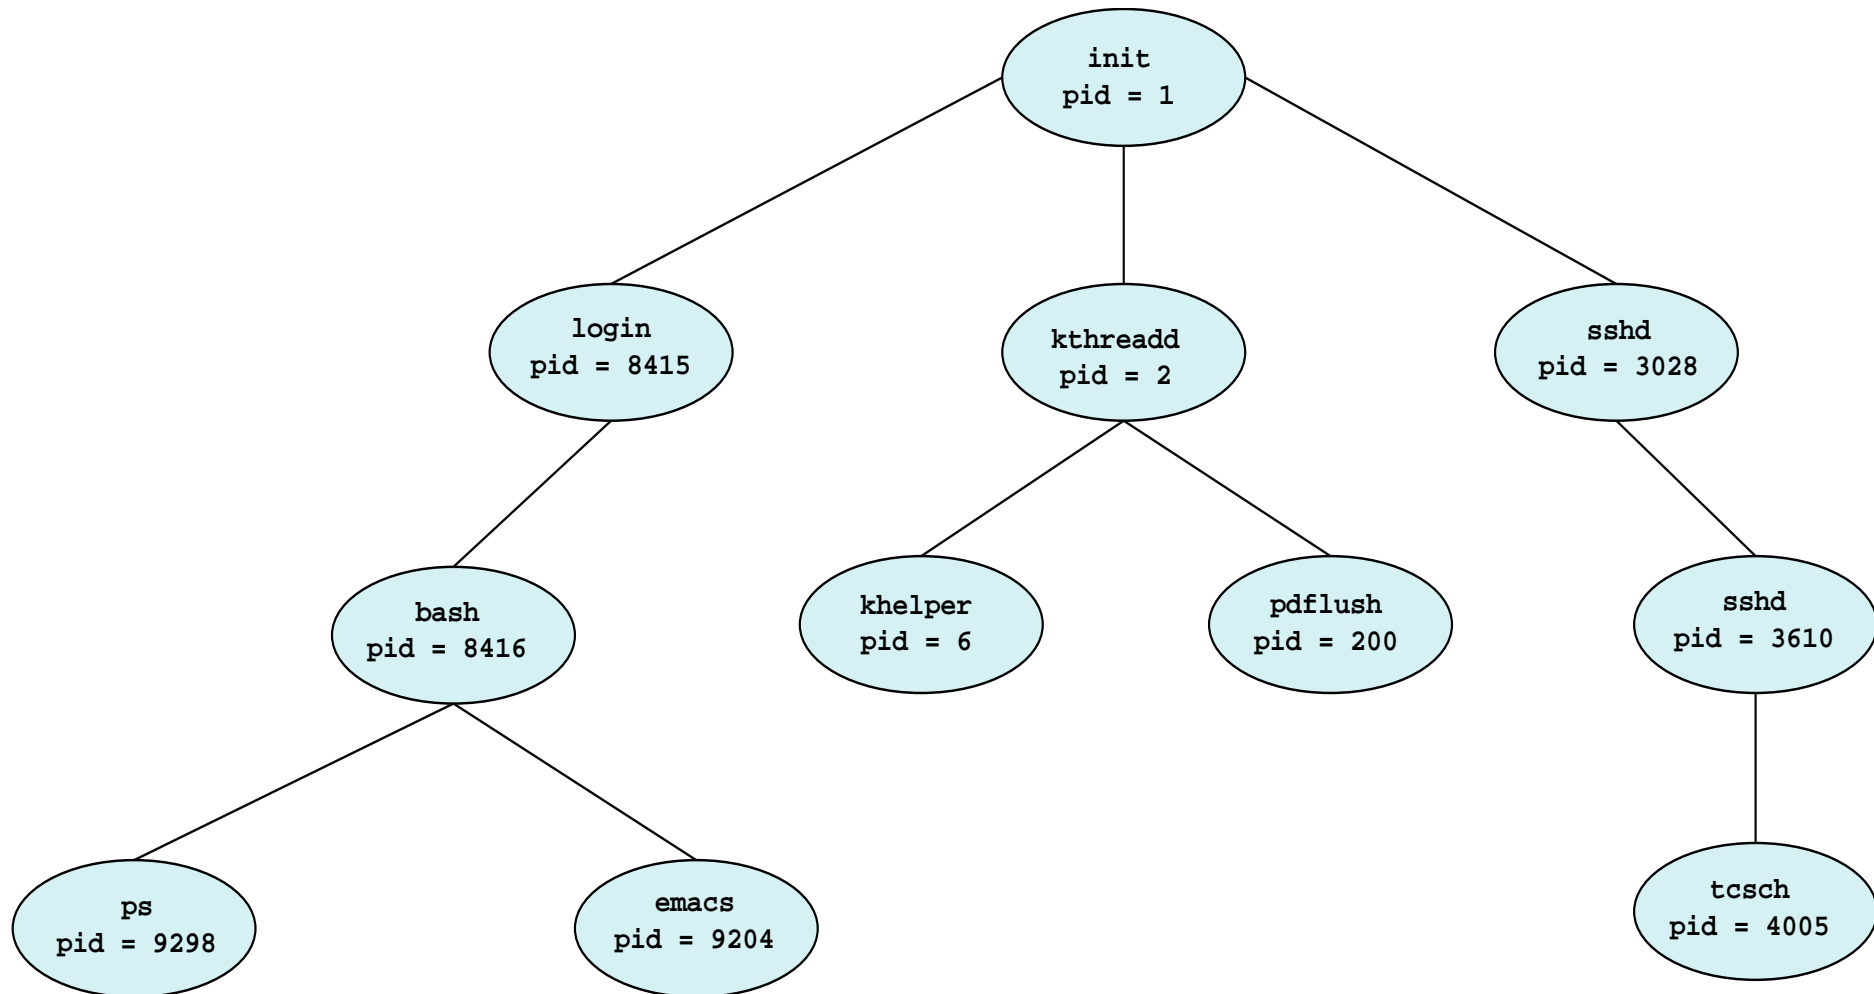


Simplified Unix/Linux Shell

```
int main()
{
    while (1) {
        char *cmd = read_command();
        int pid;
        if ((pid = fork()) == 0) {
            /* Manipulate stdin/stdout/stderr for
               pipes and redirections, etc. */
            exec(cmd);
            panic("exec failed!");
        } else {
            wait (pid);
        }
    }
}
```



A Process Tree in Linux



Process Creation/Execution: Windows

```
BOOL CreateProcess (char *prog, char *args, ...)
```

■ CreateProcess()

- ✓ Creates and initializes a new PCB
- ✓ Creates and initializes a new address space
- ✓ Loads the program specified by “prog” into the address space
- ✓ Copies “args” into memory allocated in address space
- ✓ Initializes the hardware context to start execution at main
- ✓ Places the PCB on the ready queue



Process Termination

■ Normal termination

- ✓ return from `main()`
- ✓ calling `exit()`
- ✓ calling `_exit()`

■ Abnormal termination


- ✓ calling `abort()`
- ✓ terminated by a signal

■ Wait for termination of a child process

- ✓ calling `wait()`
- ✓ If no parent waiting (did not invoke `wait()`), process is a **zombie**
- ✓ If parent terminated without invoking `wait`, process is an **orphan**



Start and Termination of a C Program



Multiprocess in Application Program

- Google Chrome Browser is multiprocess with 3 different types of processes:
 - ✓ Browser process manages user interface, disk and network I/O
 - ✓ Renderer process renders web pages, deals with HTML, Javascript
 - A new renderer created for each website opened
 - Runs in sandbox restricting disk and network I/O, minimizing effect of security exploits
 - ✓ Plug-in process for each type of plug-in



Multiprocess in Mobile Systems

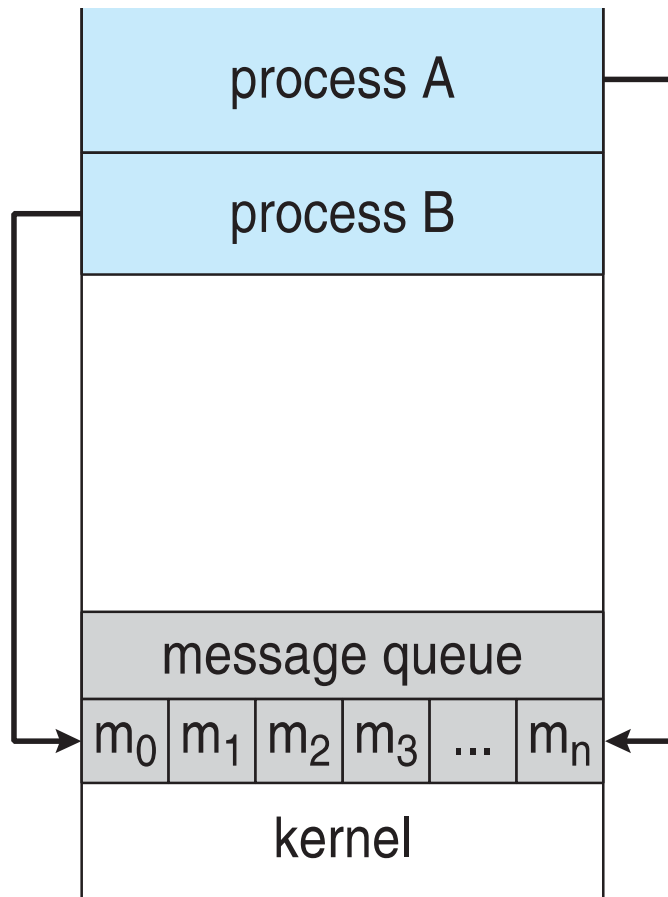
- Some mobile systems (e.g., early version of iOS) allow only one process to run, others suspended
- Due to screen real estate, user interface limits iOS provides for a
 - ✓ Single **foreground** process- controlled via user interface
 - ✓ Multiple **background** processes— in memory, running, but not on the display, and with limits
 - ✓ Limits include single, short task, receiving notification of events, specific long-running tasks like audio playback
- Android runs foreground and background, with fewer limits
 - ✓ Background process uses a **service** to perform tasks
 - ✓ Service can keep running even if background process is suspended
 - ✓ Service has no user interface, small memory use




Inter-Process Communication (IPC)

■ Communication models

- ✓ (a) message passing vs. (b) shared memory



(a)



(b)




Inter-Process Communication (IPC)

■ Cooperating processes

- ✓ Example: Bounded buffer problem (Producer-Consumer problem)

Producer

```
void producer(data)
{
    while (count==N) ;
    buffer[in] = data;
    in = (in+1) % N;
    count++;
}
```



Consumer

```
void consumer(data)
{
    while (count==0) ;
    data = buffer[out];
    out = (out+1) % N;
    count--;
}
```



Inter-Process Communication (IPC)

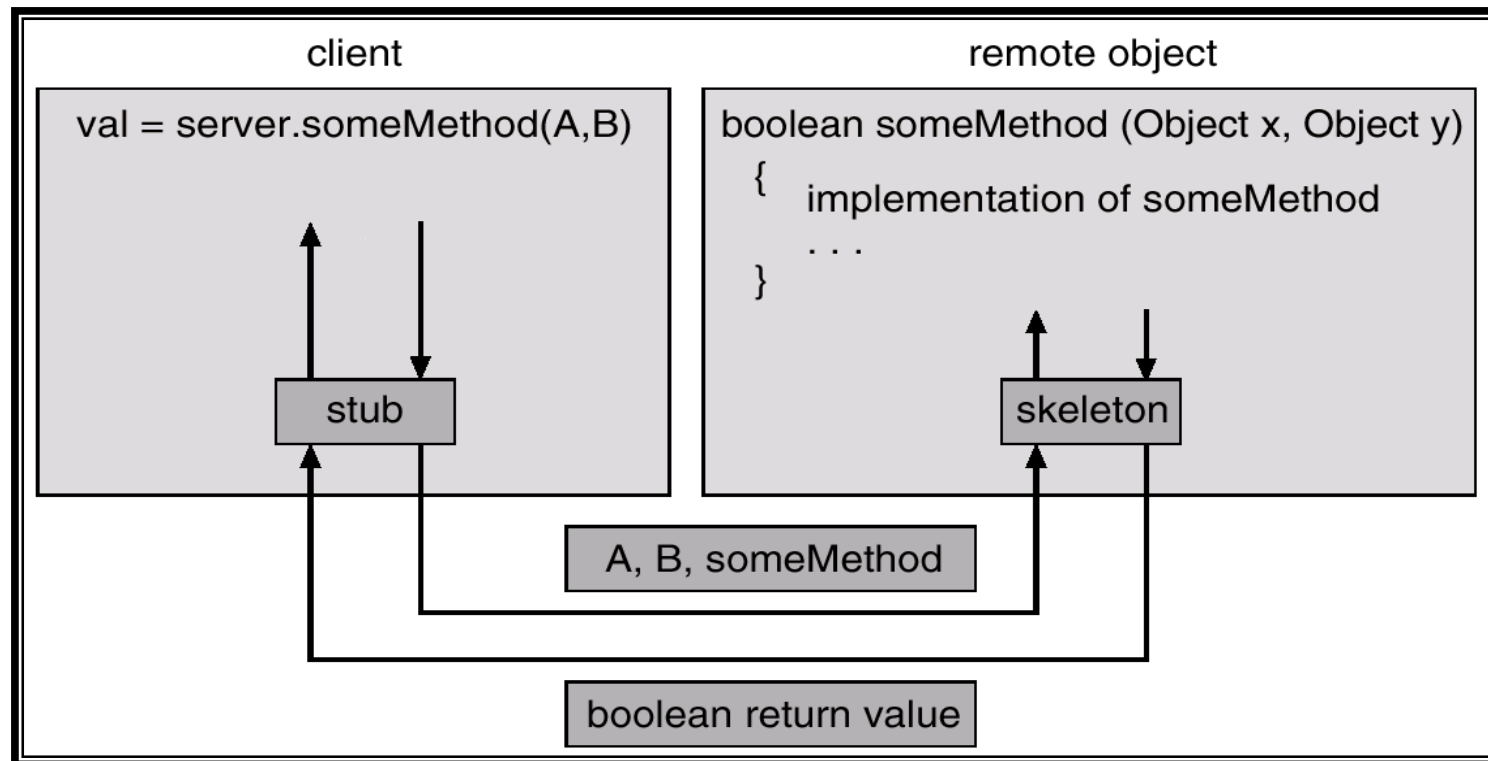
■ Unix/Linux IPC

- ✓ pipes
- ✓ FIFOs
- ✓ message queue
- ✓ shared memory
- ✓ sockets



Client-Server Communication

- Sockets
- Remote Procedure Call (RPC)
- Remote Method Invocation (RMI in Java)
- Marshalling parameters



Execution of RPC

