



Chap. 14) File System Implementation

Chap. 15) File System Internals

경희대학교 컴퓨터공학과

조진성

Overview

User's view on file systems:

- ✓ How files are named?
- ✓ What operations are allowed on them?
- ✓ What the directory tree looks like?

Implementer's view on file systems:

- ✓ How files and directories are stored?
- ✓ How disk space is managed?
- ✓ How to make everything work efficiently and reliably?



File System Implementation

In-memory structure


- ✓ In-memory partition table
- ✓ In-memory directory structure
- ✓ System-wide open file table
- ✓ Per-process open file table

On-disk structure


- ✓ Boot control block
 - Boot block(UFS) or Boot sector(NTFS)
- ✓ Volume control block
 - Super block(UFS) or Master file table(NTFS)
- ✓ Directory structure
- ✓ File control block (FCB)
 - i-node(UFS) or in master file table(NTFS)



In-Memory Structure




(a)



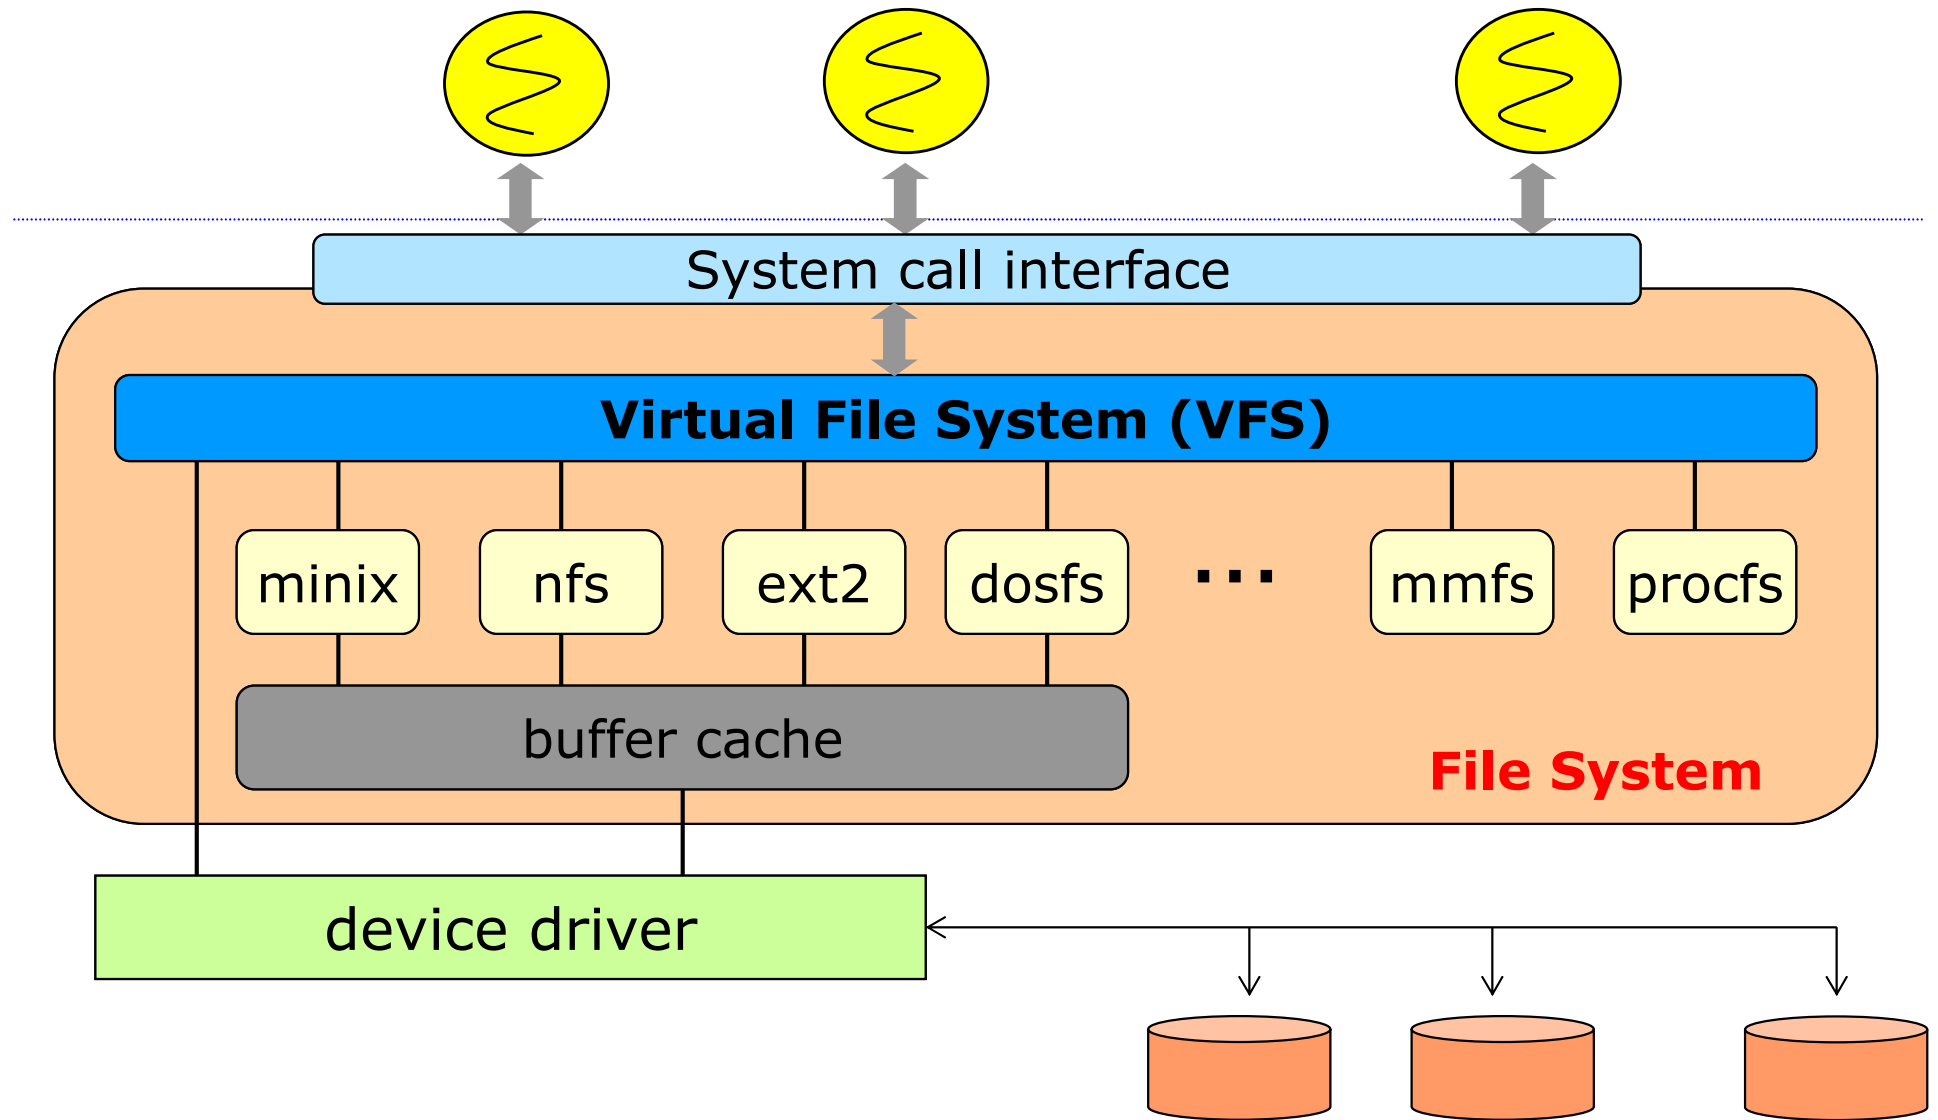
(b)




In-Memory Structure



Virtual File System



Layered File System



CD-ROM

- ✓ ISO 9660

Unix

- ✓ UFS, FFS

Windows

- ✓ FAT, FAT32, NTFS

Linux

- ✓ Over 130 different file systems
- ✓ Extended file system, ext2/3/4
- ✓ Distributed file system

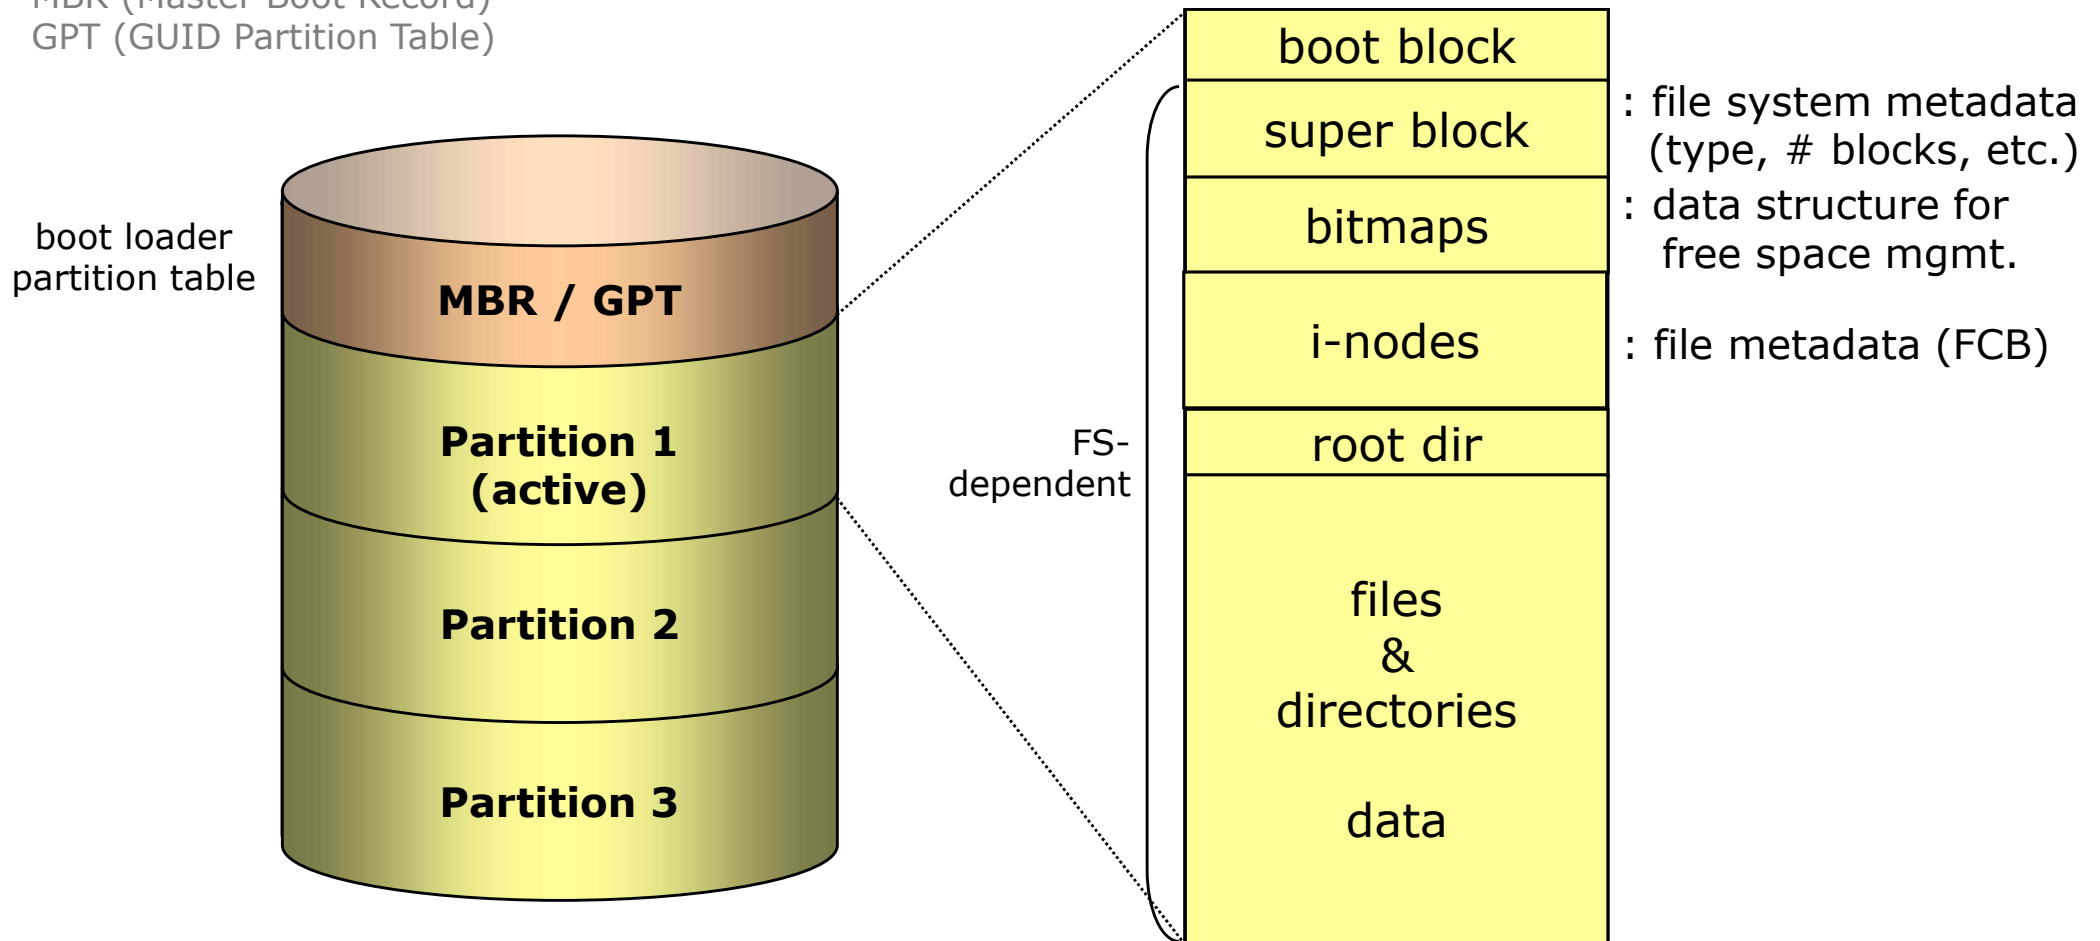
New ones

- ✓ ZFS, GoogleFS, Oracle ASM, FUSE, ...




On-Disk Structure


BIOS (Basic Input/Output System)
UEFI (Unified Extensible Firmware Interface)
MBR (Master Boot Record)
GPT (GUID Partition Table)



Disk Block



**cylinder #, surface #,
track #, sector #**




disk blocks

disk block #

Free-Space Management

Bit vector or bit map (n blocks)



$$\text{bit}[i] = \begin{cases} 1 \Rightarrow \text{block}[i] \text{ free} \\ 0 \Rightarrow \text{block}[i] \text{ occupied} \end{cases}$$

Bit map requires extra space. Example:

block size = 2^{12} bytes

disk size = 2^{30} bytes (1 gigabyte)

$n = 2^{30}/2^{12} = 2^{18}$ bits (or 32K bytes)

Easy to get contiguous files



Free-Space Management


Linked list (free list)

Grouping

Counting

Space maps

✓ ZFS



A Typical File Control Block

i-node (UFS) or in master file table (NTFS)

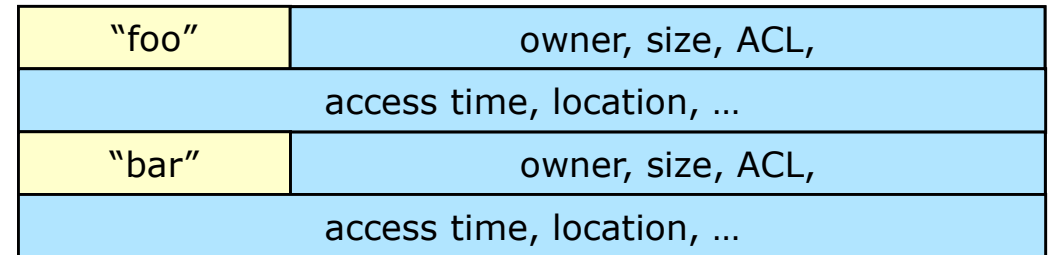
file permissions
file dates (create, access, write)
file owner, group, ACL
file size
file data blocks or pointers to file data blocks



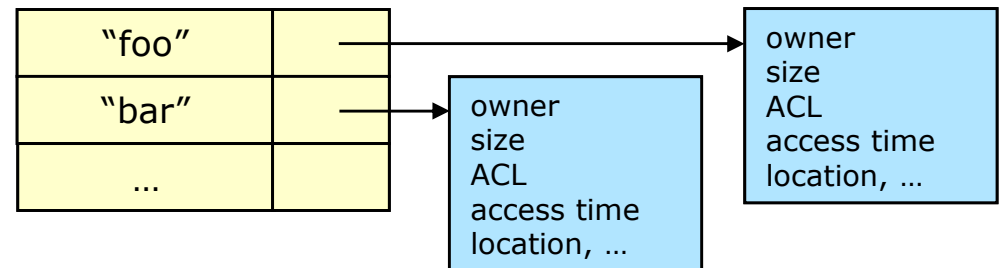
Directory Implementation

The location of metadata (FCB)


✓ In the directory entry



✓ In the separate data structure (e.g., i-node)



✓ A hybrid approach




Allocation Methods

An allocation method refers to how disk blocks are allocated for files

- ✓ Contiguous allocation
- ✓ Linked allocation
- ✓ Indexed allocation



Contiguous Allocation



directory

file	start	length
count	0	2
tr	14	3
mail	19	6
list	28	4
f	6	2



Contiguous Allocation

Advantages

- ✓ The number of disk seeks is minimal
- ✓ Directory entries can be simple:
<file name, starting disk block, length, etc.>

Disadvantages

- ✓ Requires a dynamic storage allocation: First / best fit
- ✓ External fragmentation: may require a compaction
- ✓ The file size is hard to predict and varying over time

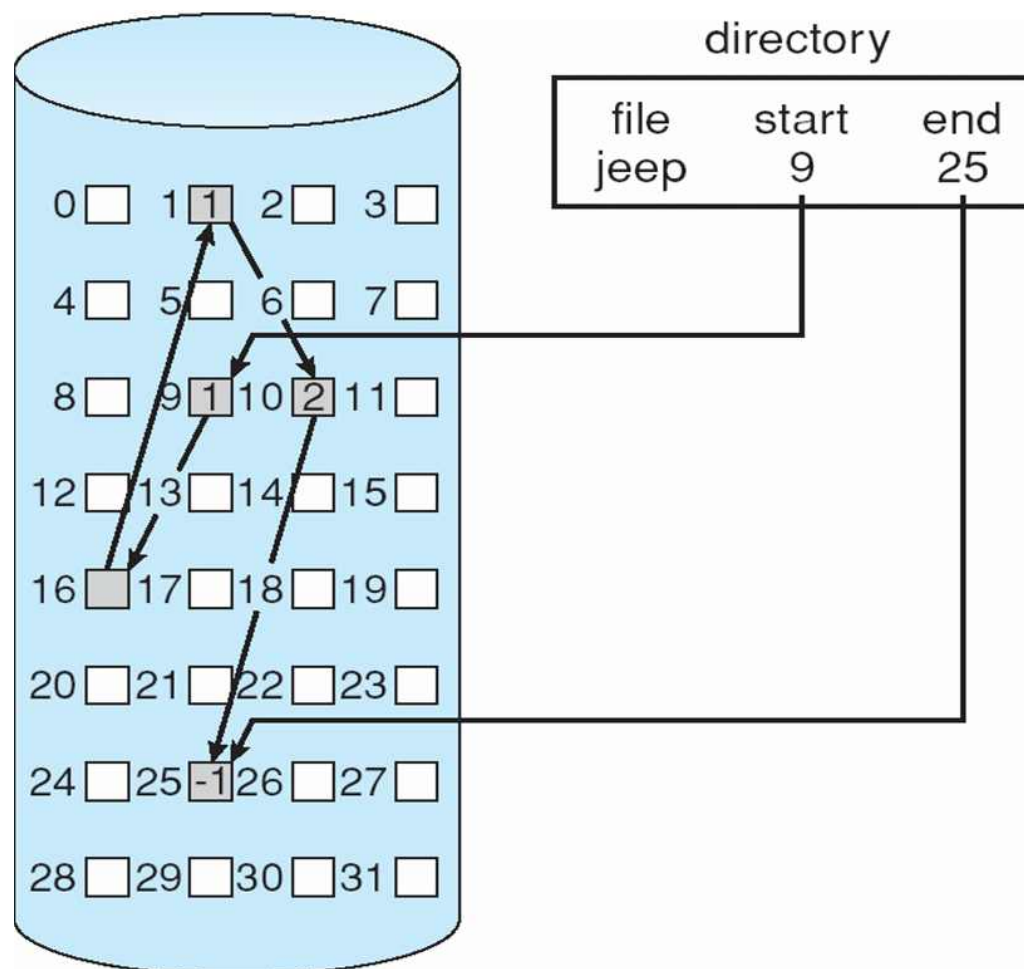

Feasible and widely used for CD-ROMS

- ✓ All the file sizes are known in advance
- ✓ Files will never change during subsequent use



Linked Allocation

block =



Linked Allocation

Advantages

- ✓ Directory entries are simple:
<file name, starting block, ending block, etc.>
- ✓ No external fragmentation
 - the disk blocks may be scattered anywhere on the disk
- ✓ A file can continue to grow as long as free blocks are available


Disadvantages

- ✓ It can be used only for sequentially accessed files
- ✓ Space overhead for maintaining pointers to the next disk block
- ✓ The amount of data storage in a block is no longer a power of two because the pointer takes up a few bytes
- ✓ Fragile: a pointer can be lost or damaged




Linked Allocation

File-Allocation Table (FAT)



Indexed Allocation



Indexed Allocation

Advantages

- ✓ Supports direct access, without suffering from external fragmentation
- ✓ I-node need only be in memory when the corresponding file is open

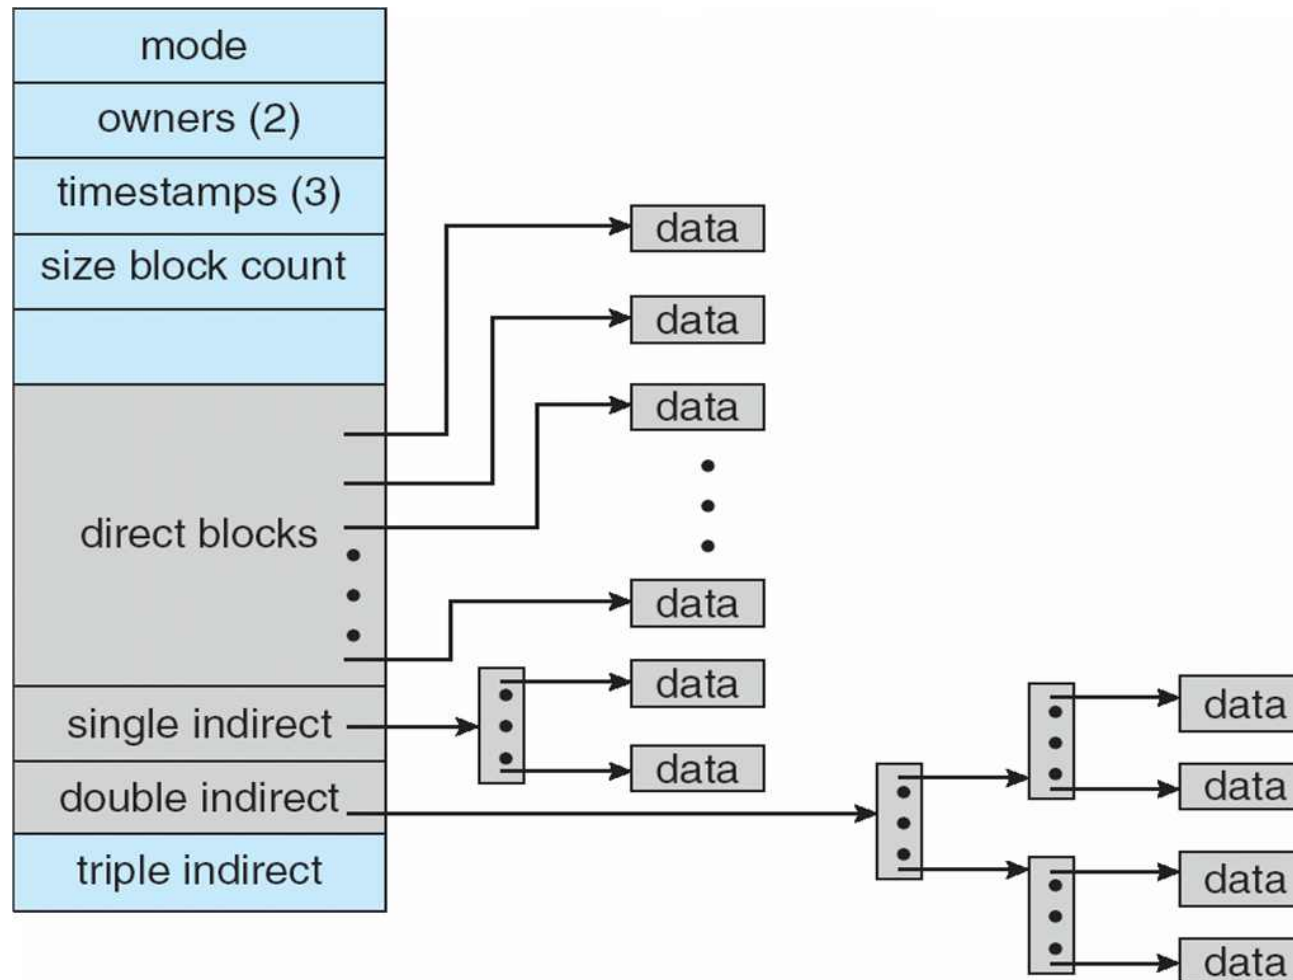
Disadvantages

- ✓ Space overhead for indexes:
 - (1) Linked scheme: link several index blocks
 - (2) Multilevel index blocks
 - (3) Combined scheme: UNIX UFS
 - 12 direct blocks, single indirect block, double indirect block, triple indirect block



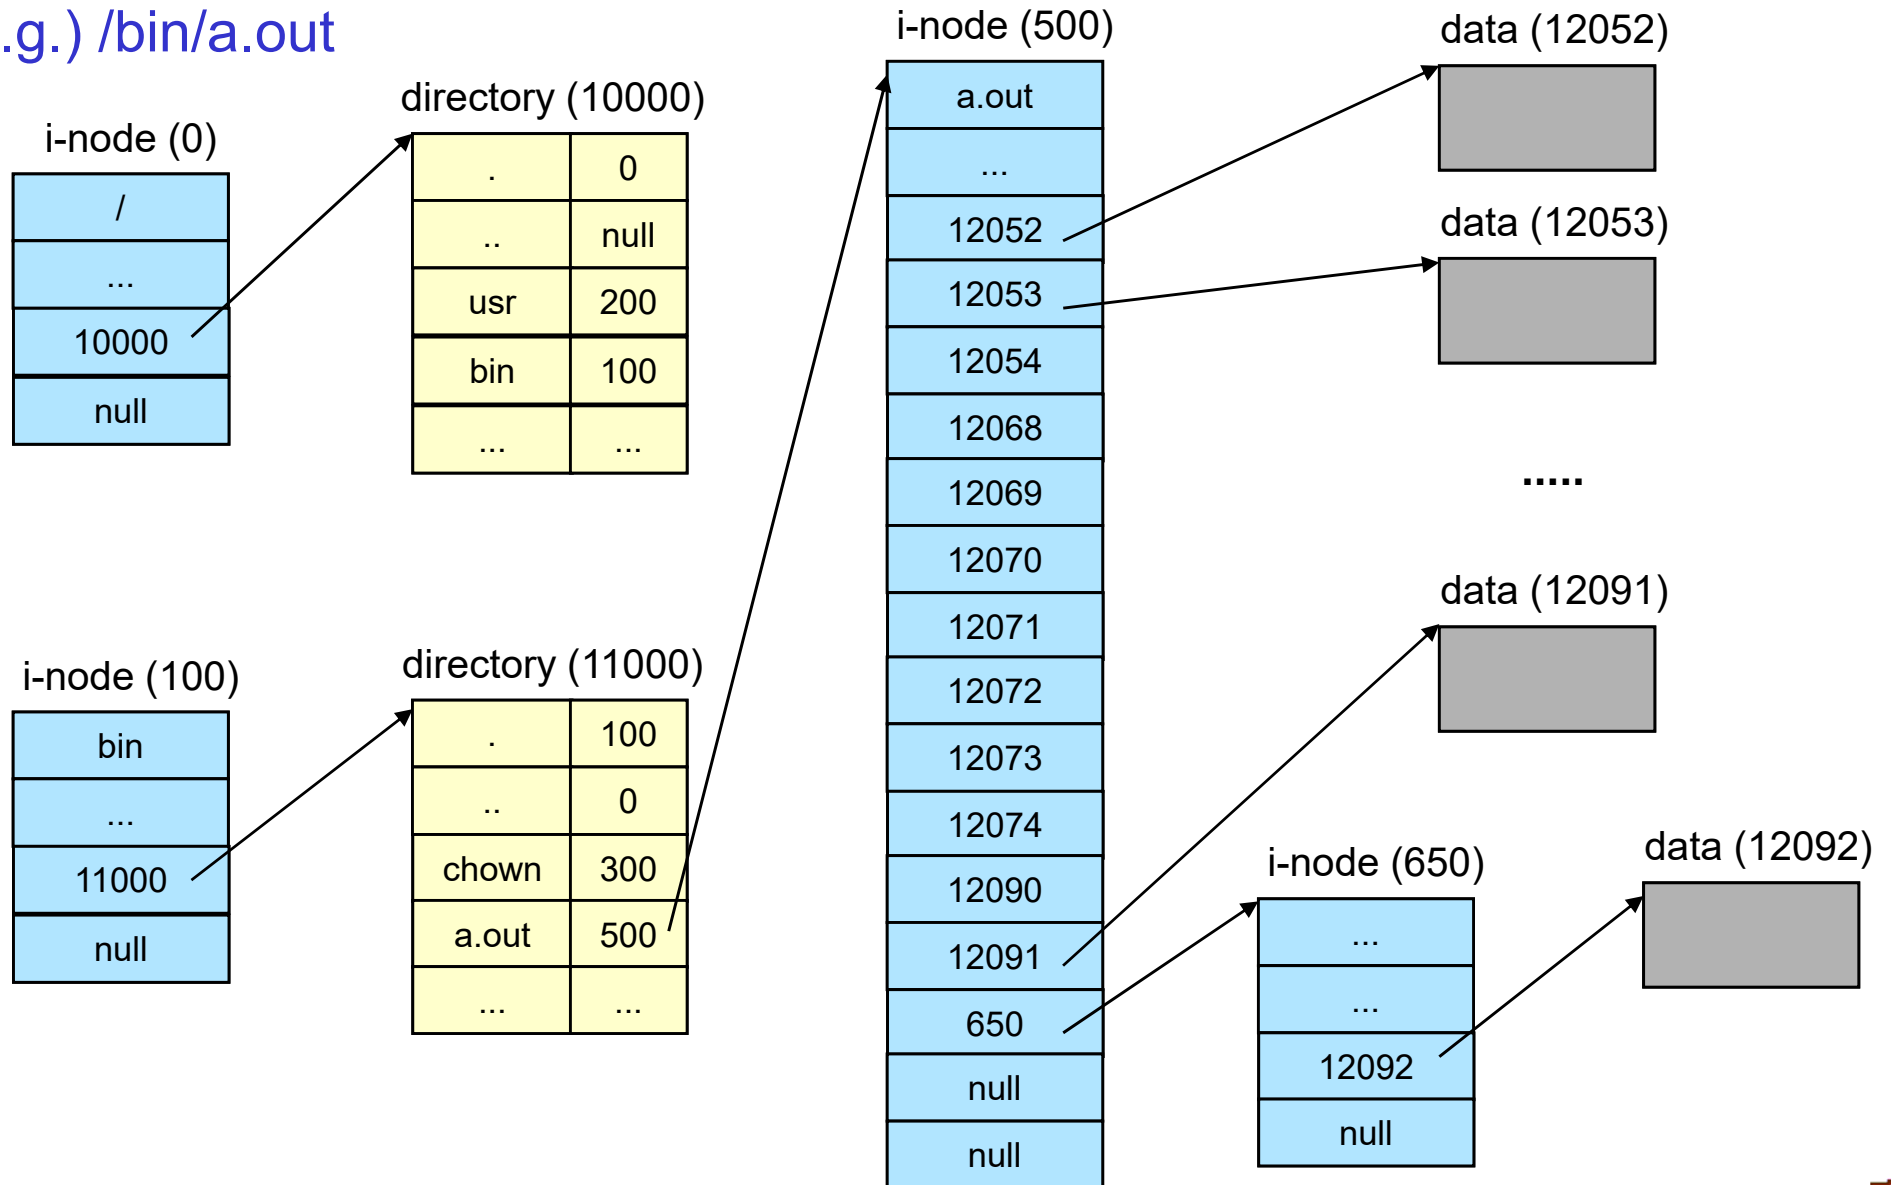
Indexed Allocation

Combined scheme: UFS (4K bytes per block, 32 bits addresses)



UNIX File System (UFS) Structure


E.g.) /bin/a.out



Block Size Performance vs. Efficiency

Block size

- ✓ Disk block size vs. file system block size
- ✓ The median file size in UNIX is about 1KB



Read-Ahead

File system predicts that the process will request next block

- ✓ File system goes ahead and requests it from the disk
- ✓ This can happen while the process is computing on previous block, overlapping I/O with execution
- ✓ When the process requests block, it will be in cache

Compliments the disk cache, which also is doing read ahead

Very effective for sequentially accessed files

File systems try to prevent blocks from being scattered across the disk during allocation or by restructuring periodically

Cf) Free-behind



Buffer Cache

Applications exhibit significant locality for reading and writing files

Idea: cache file blocks in memory to capture locality in buffer cache (or disk cache)

- ✓ Cache is system wide, used and shared by all processes
- ✓ Reading from the cache makes a disk perform like memory
- ✓ Even a 4MB cache can be very effective

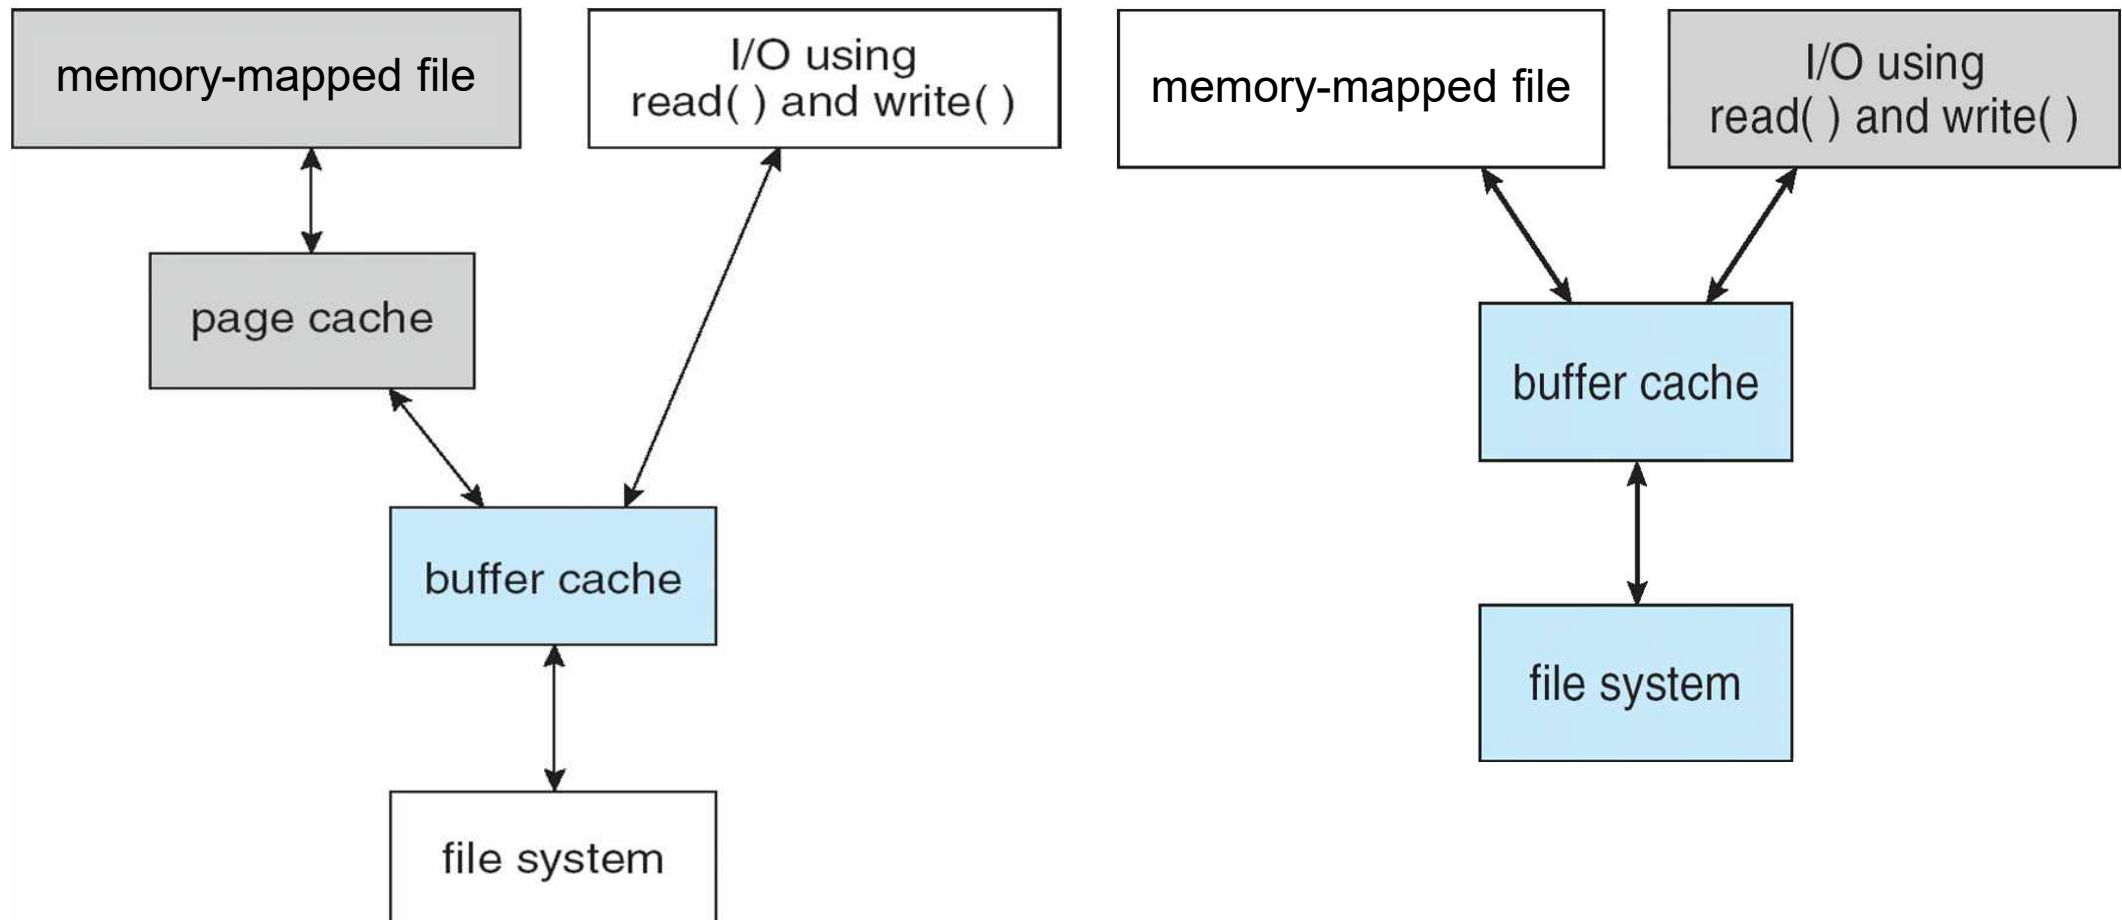
Issues

- ✓ The buffer cache competes with VM
- ✓ Like VM, it has limited size
- ✓ Need replacement algorithms again
(References are relatively infrequent, so it is feasible to keep all the blocks in exact LRU order)



Unified Buffer Cache

Page cache for demand paging (including memory-mapped file)
Buffer cache for ordinary file system I/O



Caching Writes

Synchronous writes are very slow

Asynchronous writes (or write-behind, write-back)

- ✓ Maintain a queue of uncommitted blocks
- ✓ Periodically flush the queue to disk
- ✓ **Unreliable**: metadata requires synchronous writes (with small files, most writes are to metadata)



File system consistency

- ✓ File system can be left in an inconsistent state if cached blocks are not written out due to the system crash
- ✓ It is especially critical if some of those blocks are i-node blocks, directory blocks, or blocks containing the free list
- ✓ Most systems have a utility program that checks file system consistency
 - Windows: scandisk
 - UNIX: fsck



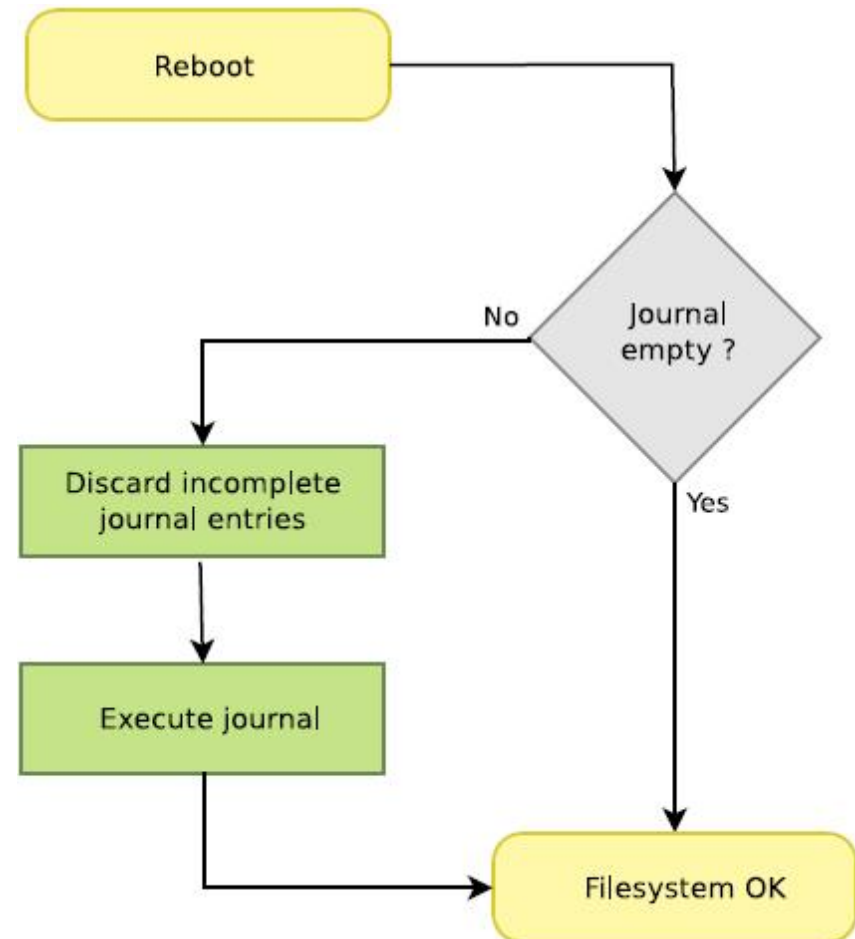
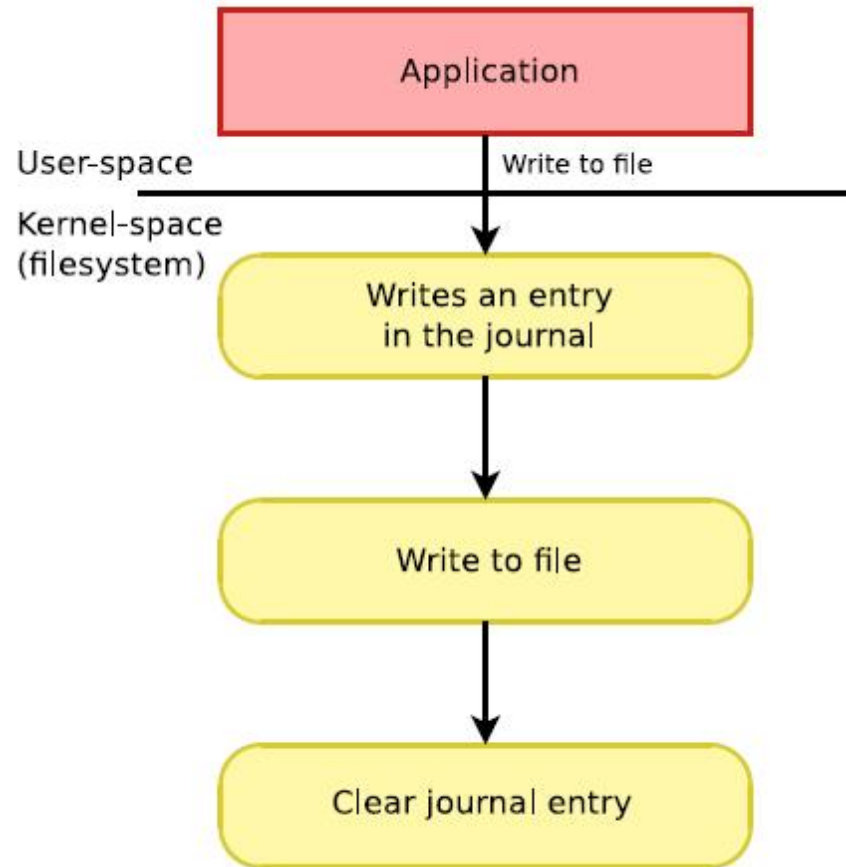
Log Structured File Systems

Journaling file systems

- ✓ Fsck'ing takes a long time, which makes the file system restart slow in the event of system crash
- ✓ Record a log, or journal, of changes made to files and directories to a separate location (preferably a separate disk)
- ✓ If a crash occurs, the journal can be used to undo any partially completed tasks that would leave the file system in an inconsistent state
- ✓ IBM JFS for AIX, Linux
 - Veritas VxFS for Solaris, HP-UX, Unixware, etc.
 - SGI XFS for IRIX, Linux
 - Reiserfs, ext3 for Linux
 - NTFS for Windows



Journaling File Systems



Remote File Systems

Network File Systems (NFS)

